

The model encodes privacy far better than its own upcoming behaviour

Identical extraction point (layer 20, final prompt token), probe family, and CV scheme; both probes see the same words. Knowing-vs-doing gap: +0.283 AUC.

