



Pre-response activations predict upcoming disclosure STRATEGY beyond the scenario wording

analysis population, n=216 · circle = layer-20 activation probe, open square = frozen sentence-encoder baseline, grey tick = TF-IDF · filled = survives Holm correction · provisional single-judge labels

(a) activation vs scenario-text channels



(b) privileged increment



Δ is measured against the STRONGER text baseline per contrast (the encoder on all six). A positive Δ means the model's processing makes behaviour-relevant information linearly extractable where text classifiers cannot extract it — not that the activation holds information absent from the prompt, which is impossible since it is a deterministic function of it. Δ is relative to this baseline family; a stronger encoder can only shrink it.